

AVL Trees


Height of an AVL Tree
Insertion and restructuring
Removal and restructuring
Costs

AVL tree, named after the initials of its inventors: **A**del'son-**V**el'skii and **L**andis

1

AVL Tree

- AVL trees are **balanced**.
- An AVL Tree is a **binary search tree** such that for every internal node v of T , the **heights of the children of v can differ by at most 1**.




An example of an AVL tree where the heights are shown next to the nodes:

2

Balancing Factor

$\text{height}(\text{right s.a.}) - \text{height}(\text{left s.a.})$


$\in \{-1, 0, 1\}$ for AVL tree



3

Height of an AVL Tree


- **Proposition:** The **height** of an AVL tree T storing n keys is $O(\log n)$.
- **Justification:** The easiest way to approach this problem is to find $n(h)$: the **minimum number of internal nodes of an AVL tree of height h** .
- We see that $n(1) = 1$ and $n(2) = 2$



4

$n(h)$: the *minimum number of internal nodes* of an AVL tree of height h .

For $n \geq 3$, an AVL tree of height h contains the root node, one AVL subtree of height $h-1$ and the other AVL subtree of height $h-2$.




i.e. $n(h) = 1 + n(h-1) + n(h-2)$

5

Height of an AVL Tree

$$n(h) = 1 + n(h-1) + n(h-2)$$



But: $n(h-1) > n(h-2)$,

that is: $n(h) > 1 + n(h-2) + n(h-2)$
which is $> 2n(h-2)$

so $n(h) > 2n(h-2)$

6

Height of an AVL Tree

So, now we know: $n(h) > 2n(h-2)$
but then also: $n(h-2) > 2n(h-4)$ } $n(h) > 2n(h-4) = 2 \cdot 2n(h-4)$

$$n(h) > 4n(h-4)$$

but then also: $n(h-4) > 2n(h-6)$

$$n(h) > 8n(h-4)$$

We can continue:

$$n(h) > 2n(h-2)$$

$$n(h) > 4n(h-4)$$

...

$$n(h) > 2^i n(h-2i)$$

7

$$n(h) > 2^i n(h-2i)$$

with

$$n(1) = 1$$

$$n(2) = 2$$

$$h-2i = 1$$

$$\text{for } i = h/2 - 1$$

$$n(h) > 2^{h/2 - 1} n(1)$$

$$n(h) > 2^{h/2 - 1}$$

$$\log n(h) > \log 2^{h/2 - 1}$$

$$\log n(h) > h/2 - 1$$

$$h < 2 \log n(h) + 2$$

which means that h is $O(\log n)$


8

Insertion

- A binary search tree T is called **balanced** if for every node v , the height of v 's children differ by at most one.
- Inserting a node into an AVL tree involves performing an **expandExternal(w)** on T , which changes the heights of some of the nodes in T .
- If an insertion causes T to become **unbalanced** we have to rebalance...

9

Insertion



10

Rebalancing after insertion

We are going to identify 3 nodes which form a grandparent, parent, child triplet and the 4 subtrees attached to them. We will rearrange these elements to create a new balanced tree.

11

Rebalancing

Step 1: Trace the path back from the point of insertion to the first node whose **grandparent is unbalanced**. Label this node x , its parent y , and grandparent z .

Examples




12

Rebalancing

Step 2: These nodes will have 4 subtrees connected to them. Label them T_1, T_2, T_3, T_4 from left to right.

Examples




13

Rebalancing

Step 3: Rename x, y, z to a, b, c according to their inorder traversal i.e. if $y < x < z$ then label y 'a', x 'b' and z 'c'.


Example



14

Rebalancing

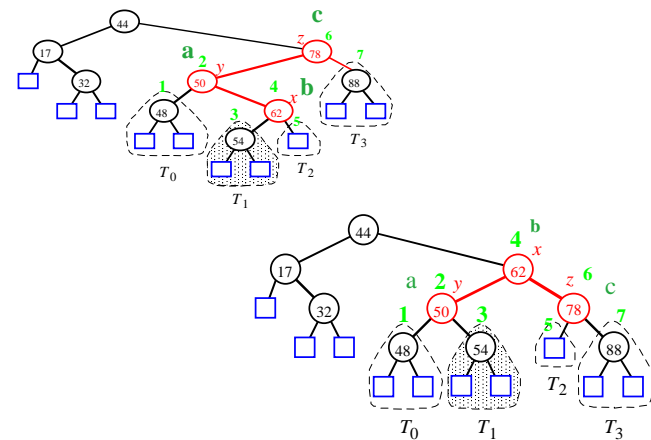
Step 4: Replace the tree rooted at z with the following tree:



Rebalance done!

15

Example



16

Does this really work?


We need to see that the new tree is :

- a) A Binary search tree - the inorder traversal of our new tree should be the same as that of the old tree
- b) Balanced: have we fixed the problem?

We consider 2 types of examples

17


Example 1



Inorder: T1 z T2 y T3 x T4

18

Example 2



Inorder: T1 y T2 x T3 z T4

19

An Observation...

Notice that in both cases, the new tree rooted at b has the same height as the old tree rooted at z had before insertion.

So.. once we have done one rebalancing act, we are done.

20

rebalance (v)

```

x <- v; Y <- x.parent; z <- y.parent
while (z.isBalanced and not(z.isRoot))
  x <- y; y <- z; z <- z.parent
if (not z.isBalanced)
  if (x = y.left) { x <- y }
  if (y = z.left) { x <- y }
  a <- x; b <- y; c <- z;
  T2 <- x.right; T3 <- y.right;
  else { z <- x }
  a <- z; b <- x; c <- y;
  T2 <- x.left; T3 <- x.right;
  else { y <- x }
  if (y = z.left) { y <- x }
  a <- y; b <- x; c <- z;
  T2 <- x.left; T3 <- x.right
  else { z <- y }
  a <- z; b <- y; c <- x;
  T2 <- y.left; T3 <- x.left

```

```

T1 <- a.left; T4 <- c.right
b.left <- a; b.right <- c
a.left <- T1; a.right <- T2
c.left <- T3; c.right <- T4
T1.parent <- a; T2.parent <- a
T3.parent <- b; T4.parent <- c

if (z.isRoot) then
  root <- b
  b.parent <- NULL
else if (z.isLeftChild)
  z.parent.left <- b
  else z.parent.right <- b
  b.parent <- z.parent
  a.parent <- b; c.parent <- b

```

21


Removal

- We can easily see that performing a **removeAboveExternal(w)** can cause T to become unbalanced.
- Let **z** be the **first unbalanced** node encountered while travelling up the tree from w. Also, let **y** be **the child of z with the larger height**, and let **x** be **the child of y with the larger height**.
- We can perform operation **restructure(x)** to restore balance at the subtree rooted at z.
- As this restructuring may upset the balance of another node higher in the tree, we must continue checking for balance until the root of T is reached

22

Removal (contd.)


the choice of x is not unique !!!



23

Removal (contd.)

- we could choose a different x:



24

COMPLEXITY

Searching: findElement(k):
Inserting: insertItem(k, o):
Removing: removeElement(k):

$O(\log n)$

25